



*We will be starting soon*



# Docker for Public Health Bioinformatics

Week 02 - Container Repositories and Writing Dockerfiles

**PRESENTED BY:**

Inês Mendes, PhD



# Course Introduction

# Training Workshop Resources

## Training Information, Communication, and Support

- **GitHub Repository** created to host training resources and information:
  - <https://github.com/theiagen/Mid-Atlantic-Docker4PH-2025>
- **Support contact:**
  - support@theiagen.com

# Course Agenda

## Docker for Public Health Bioinformatics

Week 2 - April 08/10, 2025

- Container Repositories and Writing Dockerfiles
- Hands-on Exercise: Writing Dockerfiles to Build Docker Images





## Goals by End of Week 2

- Learn about publicly available container registries & resources
- Understand the Dockerfile and how it is used for building docker images
- Learn best-practices for writing Dockerfiles
- Learn how to build a docker image on the command line using pre-defined Dockerfiles

**OBJECTIVE**



# Week 1 Review

# Week 1 Review

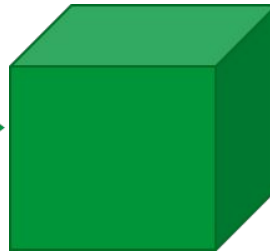
## Summary:

- **Dockerfile** is used to create the docker **image**
- Docker **image** is used to create the docker **container**
  - Container is the runnable instance of an image

Dockerfile

```
1 FROM ubuntu:xenial
2
3 # metadata
4 LABEL base.image="ubuntu:xenial"
5 LABEL version="1"
6 LABEL software="SPAdes"
7 LABEL software.version="3.13.0"
8 LABEL description="de novo DBG genome assembler"
9 LABEL website="http://cab.spbu.ru/files/release3.13.0/manual.html"
10
11 # Maintainer
12 MAINTAINER Curtis Kapsak <curtis.kapsak@state.co.us>
13
14 RUN apt-get update && apt-get install -y python \
15     wget
16
17 RUN wget http://cab.spbu.ru/files/release3.13.0/SPAdes-3.13.0-linux.tar.gz && \
18     tar -xzf SPAdes-3.13.0-linux.tar.gz && \
19     rm -r SPAdes-3.13.0-linux.tar.gz && \
20     mkdir /data
21
22 ENV PATH="${PATH}:/SPAdes-3.13.0-Linux/bin"
23 WORKDIR /data
```

Dockerfile image



docker build

docker run

Docker container





# Week 1 review

Docker Images can be built locally **or** pre-built images can be downloaded from public repositories like:

- **Docker hub:** <https://hub.docker.com/>
- **Quay.io:** <https://quay.io/>
- **GitHub container registry (GHCR):** <https://ghcr.io/>
- Cloud provider container registries:
  - GCP Artifact Registry
  - Amazon Elastic Container Registry
  - Microsoft Azure Container Registry
- Private registries are an (paid) option

# Week 1 Review

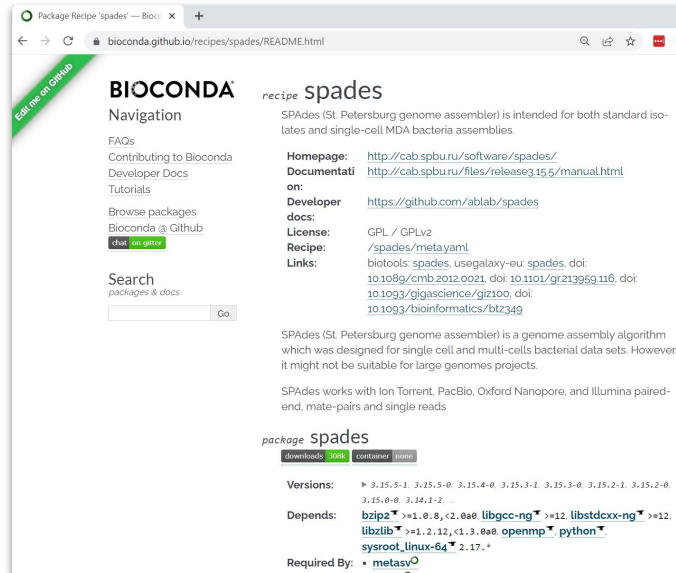
- Docker Hub: <https://hub.docker.com>
- Quay.io: <https://quay.io/>



**Red Hat**  
**Quay.io**

# Week 1 Review

- All bioconda packages are available as docker images on quay.io:
  - <https://bioconda.github.io/recipes/spades/README.html>



**BIOCONDA**

Navigation

- FAQs
- Contributing to Bioconda
- Developer Docs
- Tutorials
- Browse packages
- Bioconda @ GitHub
- [Find on GitHub](#)

Search

packages & docs

Go

**recipe spades**

SPAdes (St. Petersburg genome assembler) is intended for both standard isolates and single-cell MDA bacteria assemblies.

**Homepage:** <http://cab.spbu.ru/software/spades/>

**Documentation:** <http://cab.spbu.ru/files/release3.15.5/manual.html>

**Developer docs:** <https://github.com/ablab/spades>

**License:** GPL / GPLv2

**Recipe:** /spades/meta.yaml

**Links:** biotools: spades, usegalaxy-ou: spades, doi: 10.1089/cmb.2012.0021, doi: 10.1101/gr213959.116, doi: 10.1093/gigascience/gjz100, doi: 10.1093/bioinformatics/btz349

SPAdes (St. Petersburg genome assembler) is a genome assembly algorithm which was designed for single cell and multi-cells bacterial data sets. However, it might not be suitable for large genomes projects.

SPAdes works with Ion Torrent, PacBio, Oxford Nanopore, and Illumina paired-end, mate-pairs and single reads

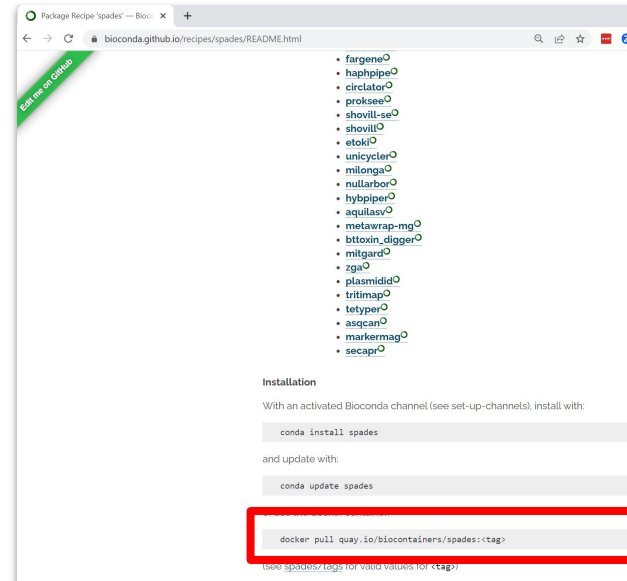
**package spades**

downloads 306k container recipe

**Versions:** 3.15.5-1, 3.15.5-0, 3.15.4-0, 3.15.3-1, 3.15.3-0, 3.15.2-1, 3.15.2-0, 3.15.0-0, 3.14.1-2, ...

**Depends:** bzp2 >=1.0.8, <2.0a0, libgcc-ng >=12, libstdc++-ng >=12, libzlib >=1.2.12, <1.3.0a0, openmp, python, sysroot\_linux-64 2.17.\*

**Required By:** metasp



**Installation**

With an activated Bioconda channel (see set-up-channels), install with:

```
conda install spades
```

and update with:

```
conda update spades
```

**docker pull quay.io/biocontainers/spades:<tag>**

(see spadesr/bugs for valid values for <tag>)



**Why Write my own  
Dockerfile?**

# Why Write my own Dockerfile?

- Not all docker images were created equally
  - Some (like biocontainers) are made by robots! 🤖
- Not all docker images work "out-of-the-box"
  - Limited-to-no testing performed with docker image
- "I cannot find a docker image for the software I want to use"
- "I want to know EXACTLY how X software was installed"
  - "Were versions pinned?"
- I want to include multiple tools in a single docker image"
  - minimap2 + samtools commands piped together:

```
minimap2 -x map-ont -a ref.fasta reads.fastq.gz | samtools sort -o out.bam
```



# Getting Started with Dockerfiles

# Review - The Dockerfile

- In order to build a docker image, you need at a minimum one file: **the Dockerfile**
  - **Dockerfile** = set of instructions used to build a docker image
  - Similar to an installation script or a **.yml** file used for making/sharing conda environments

## Format

Here is the format of the `Dockerfile` :

```
# Comment  
INSTRUCTION arguments
```

```
FROM  
ubuntu:jammy  
  
# install stuff!  
RUN install \  
    softwareA \  
    softwareB \  
    softwareC
```

# The Dockerfile

- Dockerfile instructions (**FROM**, **RUN**, **COPY**, **ENV**, etc.) will add a “layer” to the docker image
- Images are multi-layered and different images may share layers like the base image
  - **FROM** `ubuntu:focal`

## Spades Dockerfile

- <https://github.com/StaPH-B/docker-builds/blob/master/spades/3.15.5/Dockerfile>

## Dockerfile Cheat Sheet

- [https://kapeli.com/cheat\\_sheets/Dockerfile.docset/Contents/Resources/Documents/index](https://kapeli.com/cheat_sheets/Dockerfile.docset/Contents/Resources/Documents/index)

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```



# The Dockerfile - **FROM**

- Initializes a new build stage
- Required - A valid Dockerfile must start with a **FROM** instruction
  - The only instruction that can precede **FROM** is an **ARG** variable (more on this later)
- **FROM** defines the base image
  - Recommendation - choose a base image and stick with it
- Official docs:  
<https://docs.docker.com/engine/reference/builder/#from>

## FROM

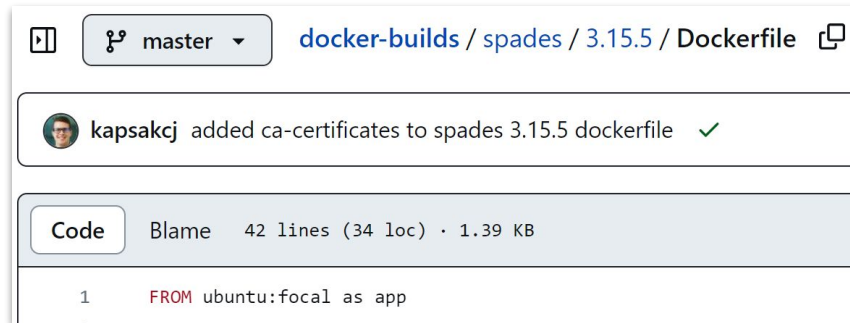
```
FROM [--platform=<platform>] <image> [AS <name>]
```

Or

```
FROM [--platform=<platform>] <image>[:<tag>] [AS <name>]
```

Or

```
FROM [--platform=<platform>] <image>[@<digest>] [AS <name>]
```



# The Dockerfile - FROM

My favorite base images

Ubuntu - [hub.docker.com/\\_/ubuntu](https://hub.docker.com/_/ubuntu)

- Familiar linux OS - has many basic linux commands installed (**ls**, **cd**, **cp**, **mv**, **ps**, etc.)
- Relatively easy to install dependencies via **apt-get**
  - [packages.ubuntu.com](http://packages.ubuntu.com) for looking up what is available via **apt-get**
- Can use **pip** for python packages & **cpan/cpanm** for perl packages

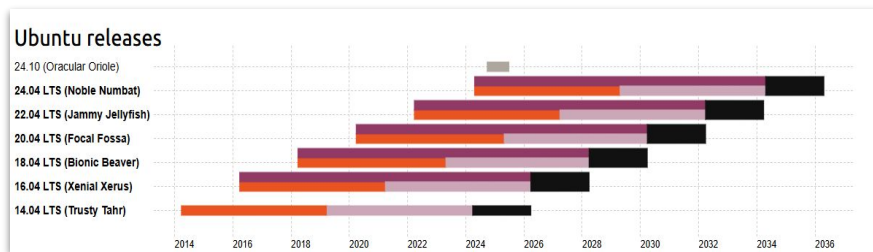
```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

# The Dockerfile - FROM

My favorite base images

Ubuntu - [hub.docker.com/\\_/ubuntu](https://hub.docker.com/_/ubuntu)

- Recommendation
  - Ubuntu 22.04 LTS (jammy)
    - **FROM** `ubuntu:jammy`
  - Ubuntu 24.04 LST (noble)
    - **FROM** `ubuntu:noble`
- Older Ubuntu releases are nearing the end of support, use something new that will be supported long term!



<https://ubuntu.com/about/release-cycle#ubuntu>

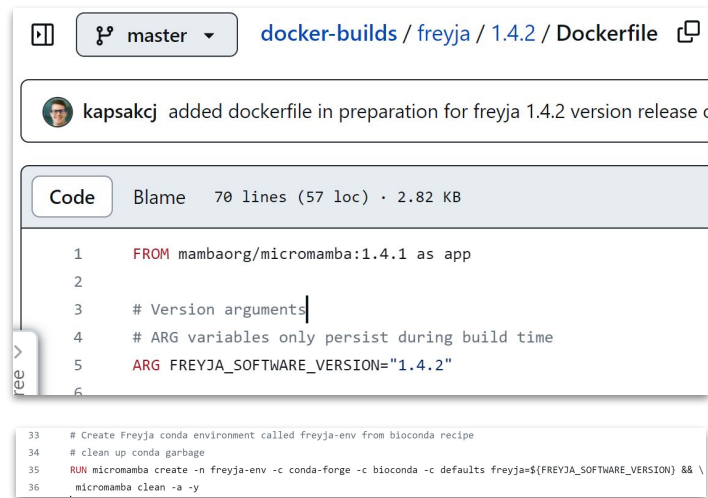
# The Dockerfile - FROM

My favorite base images

micromamba -

[hub.docker.com/r/mambaorg/micromamba](https://hub.docker.com/r/mambaorg/micromamba)

- Familiar linux OS (Debian)
- **micromamba** is preinstalled
  - **micromamba** is even more lightweight than **conda** or **miniconda**
- I use for complicated installations where I rely upon the conda package in bioconda ← not best practice!
- Also use for scenarios where I want to use a conda recipe file (.yml) for installation

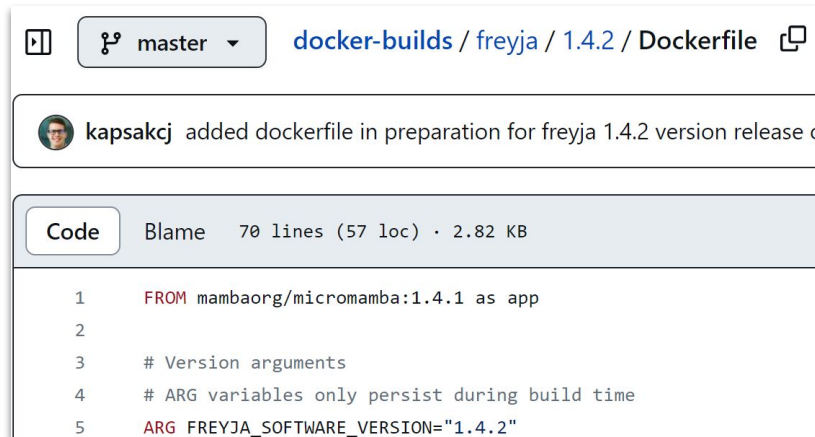


The screenshot shows a GitHub repository for 'docker-builds / freyja / 1.4.2 / Dockerfile'. The file was added by user 'kapsakcj' for preparation for freyja 1.4.2 version release. The code is 70 lines (57 loc) and 2.82 KB. The Dockerfile content is as follows:

```
1 FROM mambaorg/micromamba:1.4.1 as app
2
3 # Version arguments
4 # ARG variables only persist during build time
5 ARG FREYJA_SOFTWARE_VERSION="1.4.2"
6
33 # Create Freyja conda environment called freyja-env from bioconda recipe
34 # clean up conda garbage
35 RUN micromamba create -n freyja-env -c conda-forge -c bioconda -c defaults freyja=${FREYJA_SOFTWARE_VERSION} && \
36 micromamba clean -a -y
```

# The Dockerfile - ARG

- Sets environmental variables that are ONLY available during docker image build time
- Once image is built, all **ARG** variables are unset/removed
- Useful for specifying versions of tools to install; can make it easy to upgrade versions by only changing one line of code



```
1 FROM mambaorg/micromamba:1.4.1 as app
2
3 # Version arguments
4 # ARG variables only persist during build time
5 ARG FREYJA_SOFTWARE_VERSION="1.4.2"
```

```
33 # Create Freyja conda environment called freyja-env from bioconda recipe
34 # clean up conda garbage
35 RUN micromamba create -n freyja-env -c conda-forge -c bioconda -c defaults freyja=${FREYJA_SOFTWARE_VERSION} && \
36     micromamba clean -a -y
```

# The Dockerfile - ENV

- Sets permanent environmental variables that are available during image build and afterwards for all users
- Useful for setting the **\$PATH** variable
- [Format:](#)

**ENV** 

```
ENV <key>=<value> ...
```

```
32     # set PATH and locale settings for singularity
33     ENV LC_ALL=C.UTF-8 \
34         PATH="${PATH}:/SPADES-${SPADES_VER}-Linux/bin"
```

- Can set multiple variables in one **ENV** layer using line breaks with a backslash \
- Useful for tools that do not automatically get added to your **\$PATH** variable

# The Dockerfile - RUN

- Executes a command in a new layer
- Each **RUN** layer builds upon the previous **FROM** and **RUN** layers
- Changes made in **RUN** commands are saved in the final docker image
- Assume **/bin/sh** shell for running commands
- Docker official docs:  
<https://docs.docker.com/engine/reference/builder/#run>

## RUN

RUN has 2 forms:

- `RUN <command>` (*shell* form, the command is run in a shell, which by default is `/bin/sh -c` on Linux or `cmd /S /C` on Windows)
- `RUN ["executable", "param1", "param2"]` (*exec* form)

The `RUN` instruction will execute any commands in a new layer on top of the current image and commit the results. The resulting committed image will be used for the next step in the `Dockerfile`.

# The Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
  - **apt-get update && apt-get install** in a single **RUN** statement, done early in Dockerfile

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```



# The Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
  - **apt-get update && apt-get install** in a single **RUN** statement, done early in Dockerfile
  - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

# The Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
  - **apt-get update && apt-get install** in a single **RUN** statement, done early in Dockerfile
  - Use syntax appropriate for your build stage (shell in ubuntu). Imagine you are using the command line
  - \ (backslashes) are for line breaks (readability)

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

# The Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
  - **&&** bash operator is used to create long one-liners so that multiple commands are run sequentially and are dependent on each command running successfully

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

# The Dockerfile - RUN

- Look at a Dockerfile in-depth: SPAdes
- Tricks-of-the-trade:
  - **&&** bash operator is used to create long one-liners so that multiple commands are run sequentially and are dependent on each command running successfully
  - Assume **/bin/sh** shell for running commands, but there is a way to set **`/bin/bash`** as normal shell. Not necessary unless using bash-specific cmdline tricks

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakcj@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

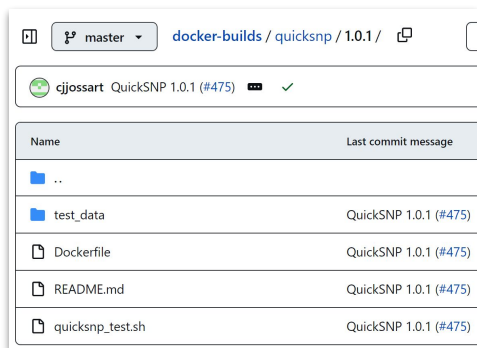
# The Dockerfile - **WORKDIR**

- Sets the working directory for any **RUN**, **CMD**, **ENTRYPOINT**, **COPY** and **ADD** instructions that follow it in the Dockerfile.
- **WORKDIR** will also persist after the image is build
- Can use multiple **WORKDIR**'s in one dockerfile, but only last **WORKDIR** will be saved in final image
- All StaPH-B/docker-builds images have the final **WORKDIR** set to **/data**

```
1 FROM ubuntu:focal as app
2
3 # to make it easier to upgrade for new versions; ARG variables only persist during docker image build time
4 ARG SPADES_VER="3.15.5"
5
6 LABEL base.image="ubuntu:focal"
7 LABEL dockerfile.version="2"
8 LABEL software="SPAdes"
9 LABEL software.version="${SPADES_VER}"
10 LABEL description="de novo DBG genome assembler"
11 LABEL website="https://github.com/ablab/spades"
12 LABEL license="https://github.com/ablab/spades/blob/v3.15.3/assembler/LICENSE"
13 LABEL maintainer="Curtis Kapsak"
14 LABEL maintainer.email="kapsakc@gmail.com"
15
16 # install dependencies; cleanup apt garbage
17 # python v3.8.10 is installed here; point 'python' to python3
18 RUN apt-get update && apt-get install --no-install-recommends -y python3 \
19     python3-distutils \
20     wget \
21     pigz \
22     ca-certificates && \
23     apt-get autoclean && rm -rf /var/lib/apt/lists/* && \
24     update-alternatives --install /usr/bin/python python /usr/bin/python3 10
25
26 # install SPAdes binary; make /data
27 RUN wget http://cab.spbu.ru/files/release${SPADES_VER}/SPAdes-${SPADES_VER}-Linux.tar.gz && \
28     tar -xzf SPAdes-${SPADES_VER}-Linux.tar.gz && \
29     rm -r SPAdes-${SPADES_VER}-Linux.tar.gz && \
30     mkdir /data
31
32 # set PATH and locale settings for singularity
33 ENV LC_ALL=C.UTF-8 \
34     PATH="${PATH}:/SPAdes-${SPADES_VER}-Linux/bin"
35
36 WORKDIR /data
37
38 # test layer
39 FROM app as test
40
41 # print version and run the supplied test flag
42 RUN spades.py --version && spades.py --test && spades.py --help
```

# The Dockerfile - COPY

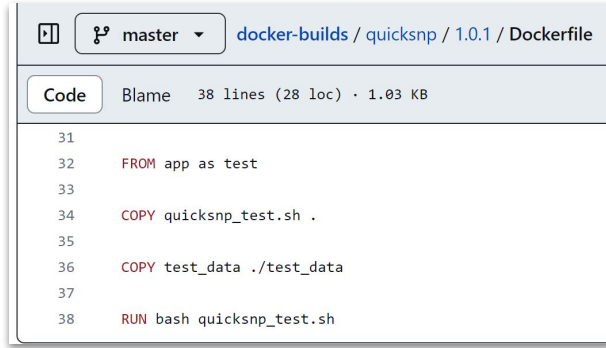
- Copies files into the image
- Files must be either:
  - located in same directory as Dockerfile (AKA the “build context”)
  - From another stage of the build process
- Useful for copying in test data, (small) databases, additional code:



master docker-builds / quicksnp / 1.0.1 /

cjossart QuickSNP 1.0.1 (#475)

Name	Last commit message
..	
test_data	QuickSNP 1.0.1 (#475)
Dockerfile	QuickSNP 1.0.1 (#475)
README.md	QuickSNP 1.0.1 (#475)
quicksnp_test.sh	QuickSNP 1.0.1 (#475)



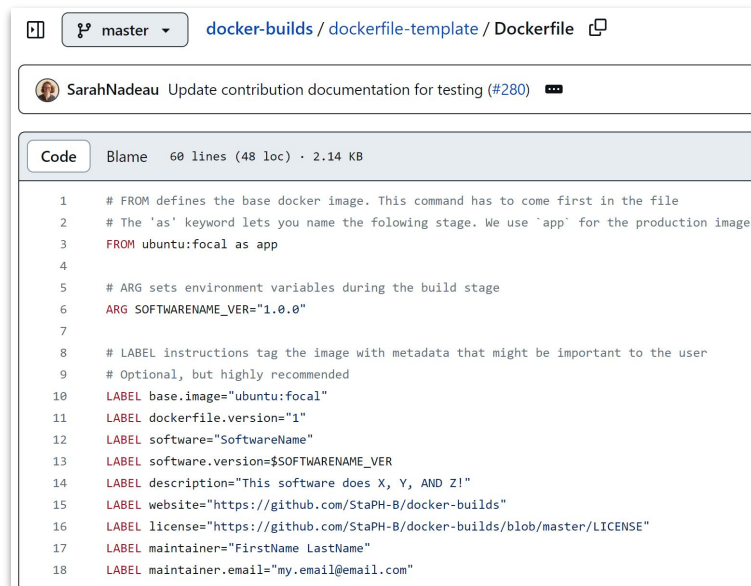
master docker-builds / quicksnp / 1.0.1 / Dockerfile

Code Blame 38 lines (28 loc) · 1.03 KB

```
31
32 FROM app as test
33
34 COPY quicksnp_test.sh .
35
36 COPY test_data ./test_data
37
38 RUN bash quicksnp_test.sh
```

# The Dockerfile - LABEL

- Optional but highly recommended!
- Allows addition of metadata to your docker image
- Generally located near top of dockerfile, after **FROM**
- See the template Dockerfile for examples:
  - <https://github.com/theiagen/Mid-Atlantic-Docker4PH-2025/blob/main/docker/example/Dockerfile>

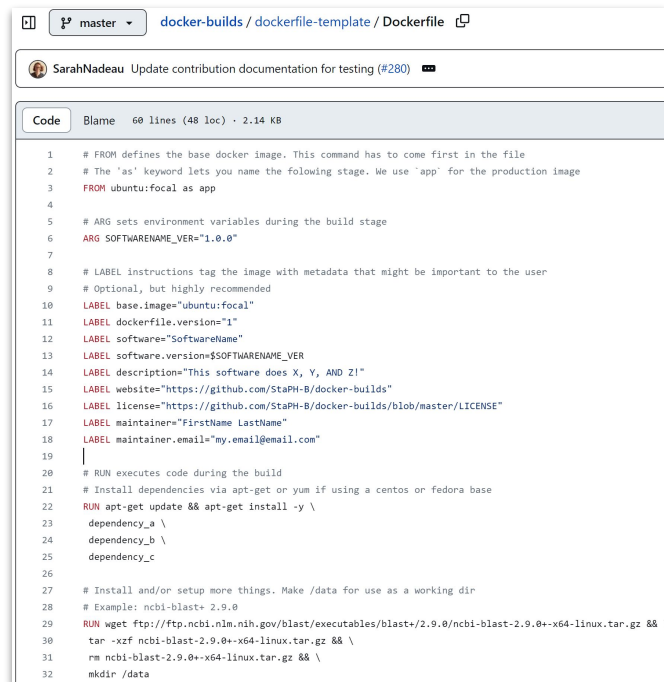


The screenshot shows a GitHub repository page for 'docker-builds / dockerfile-template / Dockerfile'. The file is 60 lines long, 48 loc, and 2.14 KB. It contains a Dockerfile template with the following content:

```
1 # FROM defines the base docker image. This command has to come first in the file
2 # The 'as' keyword lets you name the following stage. We use 'app' for the production image
3 FROM ubuntu:focal as app
4
5 # ARG sets environment variables during the build stage
6 ARG SOFTWARENAME_VER="1.0.0"
7
8 # LABEL instructions tag the image with metadata that might be important to the user
9 # Optional, but highly recommended
10 LABEL base.image="ubuntu:focal"
11 LABEL dockerfile.version="1"
12 LABEL software="SoftwareName"
13 LABEL software.version=$SOFTWARENAME_VER
14 LABEL description="This software does X, Y, AND Z!"
15 LABEL website="https://github.com/StaPH-B/docker-builds"
16 LABEL license="https://github.com/StaPH-B/docker-builds/blob/master/LICENSE"
17 LABEL maintainer="FirstName LastName"
18 LABEL maintainer.email="my.email@email.com"
```

# The Dockerfile - comments

- Also optional but highly recommended!
- Comment lines begin with #



```
1  # FROM defines the base docker image. This command has to come first in the file
2  # The 'as' keyword lets you name the following stage. We use 'app' for the production image
3  FROM ubuntu:focal as app
4
5  # ARG sets environment variables during the build stage
6  ARG SOFTWARENAME_VER="1.0.0"
7
8  # LABEL instructions tag the image with metadata that might be important to the user
9  # Optional, but highly recommended
10 LABEL base.image="ubuntu:focal"
11 LABEL dockerfile.version="1"
12 LABEL software="SoftwareName"
13 LABEL software.version=$SOFTWARENAME_VER
14 LABEL description="This software does X, Y, AND Z!"
15 LABEL website="https://github.com/StaPH-B/docker-builds"
16 LABEL license="https://github.com/StaPH-B/docker-builds/blob/master/LICENSE"
17 LABEL maintainer="FirstName LastName"
18 LABEL maintainer.email="my.email@email.com"
19 |
20 # RUN executes code during the build
21 # Install dependencies via apt-get or yum if using a centos or fedora base
22 RUN apt-get update && apt-get install -y \
23     dependency_a \
24     dependency_b \
25     dependency_c
26
27 # Install and/or setup more things. Make /data for use as a working dir
28 # Example: ncbi-blast+ 2.9.0
29 RUN wget ftp://ftp.ncbi.nlm.nih.gov/blast/executables/blast+/2.9.0/ncbi-blast-2.9.0+-x64-linux.tar.gz && \
30     tar -xzf ncbi-blast-2.9.0+-x64-linux.tar.gz && \
31     rm ncbi-blast-2.9.0+-x64-linux.tar.gz && \
32     mkdir /data
```





# Building docker images

# Docker Build

- Builds an image from a Dockerfile
  - At a minimum, requires a Dockerfile. Some Dockerfiles require other files for building (scripts, databases, etc.)

- Official docs:

<https://docs.docker.com/engine/reference/commandline/build/>

- General command structure:

```
docker build --tag <name>:<tag> <directory-with-dockerfile>
```

- example using SPAdes dockerfile:

```
docker build --tag spades:3.15.5 spades/3.15.5/
```

# Multi-Stage Docker Builds

- Docker images that use multiple **FROM** instructions
  - can use a new base image **OR**
  - can use previous stage
- Useful for “optimizing” Dockerfiles:
  - reducing number of layers
  - removing unnecessary software & files, etc. to reduce final size of docker image

```
FROM baseimage as name
RUN stuff

FROM name as name2
RUN stuff

FROM name2 as final
RUN stuff
```

# Multi-Stage Docker Builds

```
FROM ubuntu:focal as stage1
```

```
RUN stuff
```

```
FROM python:slim as stage2
```

```
RUN stuff
```

```
FROM ubuntu:jammy as finalstage
```

```
RUN stuff
```



# Hands-On Exercise

# Exercise 02: Writing Dockerfiles to Build Docker Images

## Exercise Goal:

- Access development environment via GitPod and VS Code
- [Optional] Create branch for Week 02
- Building a Pre-Existing Dockerfile
- Editing a Pre-Existing Dockerfile
- [Extra content] Pushing to a remote Docker repository (<https://hub.docker.com>)



[www.theiagen.com](http://www.theiagen.com)  
[support@theiagen.com](mailto:support@theiagen.com)